

```
In [6]: ▶ import pandas as pd
import numpy as np
```


```
In [7]: ▶ df = pd.read_csv('CarPrice.csv')
```

```
In [8]: ▶ df
```

Out[8]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewhe
0	1	3	alfa-romero giulia	gas	std	two	convertible	rv
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rv
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rv
3	4	2	audi 100 ls	gas	std	four	sedan	fv
4	5	2	audi 100ls	gas	std	four	sedan	4v
...	...	...	...	...	...	...	...	...
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rv
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rv
202	203	-1	volvo 244dl	gas	std	four	sedan	rv
203	204	-1	volvo 246	diesel	turbo	four	sedan	rv
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rv

205 rows × 26 columns



```
In [4]: ▶ ## DATA PREPROCESSING ##
```

In [5]: `df.head()`

Out[5]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd
3	4	2	audi 100 ls	gas	std	four	sedan	fwd
4	5	2	audi 100ls	gas	std	four	sedan	4wd

5 rows × 26 columns

In [6]: `df.tail()`

Out[6]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rwd
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rwd
202	203	-1	volvo 244dl	gas	std	four	sedan	rwd
203	204	-1	volvo 246	diesel	turbo	four	sedan	rwd
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rwd

5 rows × 26 columns

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null    int64
1   symboling              205 non-null    int64
2   CarName               205 non-null    object
3   fueltype              205 non-null    object
4   aspiration            205 non-null    object
5   doornumber            205 non-null    object
6   carbody               205 non-null    object
7   drivewheel           205 non-null    object
8   enginelocation        205 non-null    object
9   wheelbase             205 non-null    float64
10  carlength             205 non-null    float64
11  carwidth              205 non-null    float64
12  carheight            205 non-null    float64
13  curbweight            205 non-null    int64
14  enginetype            205 non-null    object
15  cylindernumber        205 non-null    object
16  enginesize            205 non-null    int64
17  fuelsystem            205 non-null    object
18  boreratio             205 non-null    float64
19  stroke                205 non-null    float64
20  compressionratio      205 non-null    float64
21  horsepower            205 non-null    int64
22  peakrpm              205 non-null    int64
23  citympg              205 non-null    int64
24  highwaympg           205 non-null    int64
25  price                205 non-null    float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

In [8]: `df.shape`

Out[8]: (205, 26)

In [9]: `df.columns`

Out[9]: Index(['car\_ID', 'symboling', 'CarName', 'fueltype', 'aspiration', 'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype', 'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke', 'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg', 'price'], dtype='object')

```
In [10]: df.duplicated().sum()
```

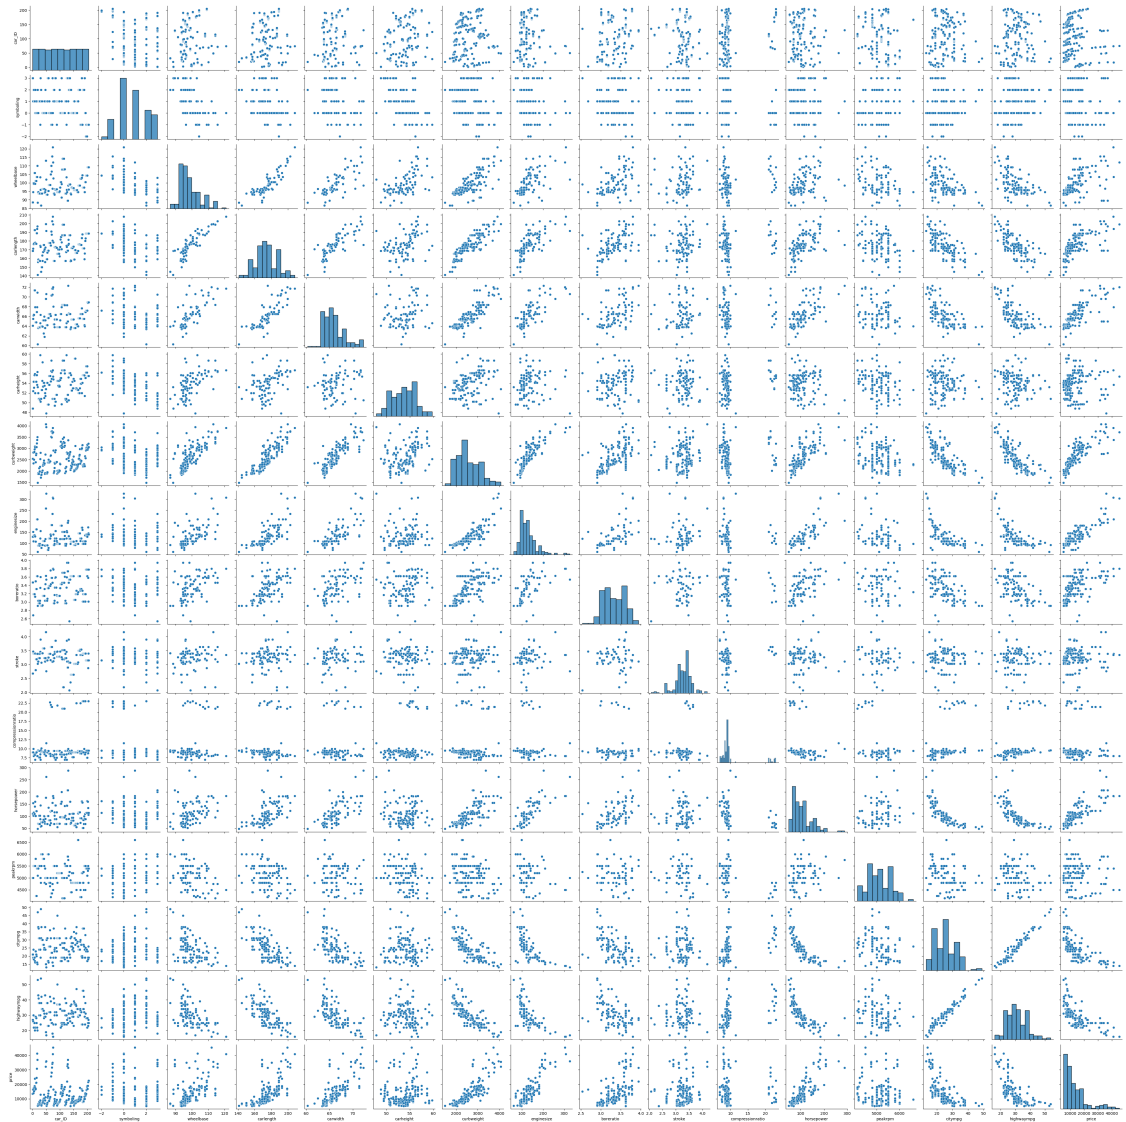
```
Out[10]: 0
```

```
In [11]: df.isnull().sum()
```


```
Out[11]: car_ID          0
          symboling      0
          CarName        0
          fueltype       0
          aspiration     0
          doornumber     0
          carbody        0
          drivewheel     0
          enginelocation 0
          wheelbase      0
          carlength      0
          carwidth       0
          carheight      0
          curbweight     0
          enginetype     0
          cylindernumber 0
          enginesize     0
          fuelsystem     0
          boreratio      0
          stroke         0
          compressionratio 0
          horsepower     0
          peakrpm        0
          citympg        0
          highwaympg     0
          price          0
          dtype: int64
```

```
In [4]: ## VISUALIZATION ##  
  
import matplotlib.pyplot as plt  
import seaborn as sns  
  
sns.pairplot(df)
```


Out[4]: <seaborn.axisgrid.PairGrid at 0x223fd082520>



```
In [6]: plt.hist(df["wheelbase"], bins=25, rwidth=0.8)
plt.show()
```



```
In [8]: plt.hist(df["carlength"], bins=10, rwidth=0.8)
plt.show()
```



In [13]:

```
print(df.fueltype.value_counts())
print(df.aspiration.value_counts())
print(df.doornumber.value_counts())
print(df.carbody.value_counts())
print(df.drivewheel.value_counts())
print(df.fuelsystem.value_counts())
print(df.cylindernumber.value_counts())
```

```
gas      185
diesel   20
Name: fueltype, dtype: int64
std      168
turbo    37
Name: aspiration, dtype: int64
four     115
two      90
Name: doornumber, dtype: int64
sedan     96
hatchback 70
wagon     25
hardtop    8
convertible 6
Name: carbody, dtype: int64
fwd      120
rwd       76
4wd        9
Name: drivewheel, dtype: int64
mpfi      94
2bbl      66
idi        20
1bbl      11
spdi        9
4bbl        3
mfi         1
spfi         1
Name: fuelsystem, dtype: int64
four      159
six        24
five       11
eight       5
two         4
three        1
twelve       1
Name: cylindernumber, dtype: int64
```

In [14]: ▶ `df['horsepower'].value_counts()`



```
Out[14]: 68      19
          70      11
          69      10
          116     9
          110     8
          95      7
          114     6
          160     6
          101     6
          62      6
          88      6
          145     5
          76      5
          97      5
          84      5
          90      5
          82      5
          102     5
          92      4
          111     4
          123     4
          86      4
          207     3
          73      3
          182     3
          121     3
          85      3
          152     3
          176     2
          94      2
          56      2
          112     2
          161     2
          184     2
          155     2
          156     2
          52      2
          100     2
          162     2
          140     1
          115     1
          134     1
          78      1
          142     1
          288     1
          143     1
          48      1
          200     1
          58      1
          55      1
          60      1
          175     1
          154     1
          72      1
          120     1
          64      1
          135     1
```

```
262      1
106      1
Name: horsepower, dtype: int64
```

```
In [15]: df['stroke'].value_counts()
```

```
Out[15]: 3.400      20
          3.230      14
          3.150      14
          3.030      14
          3.390      13
          2.640      11
          3.290       9
          3.350       9
          3.460       8
          3.110       6
          3.270       6
          3.410       6
          3.070       6
          3.580       6
          3.190       6
          3.500       6
          3.640       5
          3.520       5
          3.860       4
          3.540       4
          3.470       4
          3.255       4
          3.900       3
          2.900       3
          3.100       2
          4.170       2
          2.800       2
          2.190       2
          3.080       2
          2.680       2
          2.360       1
          3.160       1
          2.070       1
          3.210       1
          3.120       1
          2.760       1
          2.870       1
Name: stroke, dtype: int64
```

```
In [16]: df['compressionratio'].value_counts()
```

```
Out[16]: 9.00      46
          9.40      26
          8.50      14
          9.50      13
          9.30      11
          8.70       9
          8.00       8
          9.20       8
          7.00       7
          8.60       5
          21.00      5
          8.40       5
          7.50       5
          23.00      5
          9.60       5
          21.50      4
          7.60       4
          10.00      3
          22.50      3
          8.30       3
          8.80       3
          7.70       2
          8.10       2
          9.10       1
          9.31       1
          7.80       1
          9.41       1
          21.90      1
          22.00      1
          22.70      1
          10.10      1
          11.50      1
          Name: compressionratio, dtype: int64
```

```
In [17]: df['citympg'].value_counts()
```

```
Out[17]: 31      28
          19      27
          24      22
          27      14
          17      13
          26      12
          23      12
          21       8
          25       8
          30       8
          38       7
          28       7
          16       6
          37       6
          22       4
          29       3
          15       3
          20       3
          18       3
          14       2
          34       1
          35       1
          32       1
          36       1
          45       1
          13       1
          49       1
          47       1
          33       1
          Name: citympg, dtype: int64
```

```
In [18]: df['highwaympg'].value_counts()
```

```
Out[18]: 25    19
          38    17
          24    17
          30    16
          32    16
          34    14
          37    13
          28    13
          29    10
          33     9
          22     8
          31     8
          23     7
          27     5
          43     4
          42     3
          26     3
          41     3
          19     2
          39     2
          18     2
          16     2
          20     2
          36     2
          47     2
          46     2
          54     1
          17     1
          53     1
          50     1
          Name: highwaympg, dtype: int64
```

```
In [19]: ## CHANGING THE CATEGORICAL ATTRIBUTES INTO NUMERIC DATA FOR BETTER ANALYSIS
```


```
df.replace({'fueltype':{'gas':0,'diesel':1}},inplace=True)
df.replace({'aspiration':{'std':0,'turbo':1}},inplace=True)
df.replace({'doornumber':{'two':0,'four':1}},inplace=True)
df.replace({'carbody':{'convertible':0,'hatchback':1,'sedan':2, 'wagon':3})
df.replace({'drivewheel':{'rwd':0,'fwd':1,'4wd':2}},inplace=True)
df.replace({'fuelsystem':{'mpfi':0,'2bbl':1,'1bbl':2,'mfi':3, 'spf1':4, '':
```

```
In [9]: ▶ ## hot encoding ##  
df = pd.get_dummies(df, drop_first=True)  
df
```

Out[9]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize
0	1	3	88.6	168.8	64.1	48.8	2548	130
1	2	3	88.6	168.8	64.1	48.8	2548	130
2	3	1	94.5	171.2	65.5	52.4	2823	152
3	4	2	99.8	176.6	66.2	54.3	2337	109
4	5	2	99.4	176.6	66.4	54.3	2824	136
...	...	...	...	...	...	...	...	...
200	201	-1	109.1	188.8	68.9	55.5	2952	141
201	202	-1	109.1	188.8	68.8	55.5	3049	141
202	203	-1	109.1	188.8	68.9	55.5	3012	173
203	204	-1	109.1	188.8	68.9	55.5	3217	145
204	205	-1	109.1	188.8	68.9	55.5	3062	141

205 rows × 191 columns



```
In [10]: ## Splitting the data ##

x = df.drop(['price'], axis=1)
y = df['price']
print(len(x), len(y))
print(x)
print(y)
print(x.shape)
print(y.shape)
```

```

205 205
      car_ID  symboling  wheelbase  carlength  carwidth  carheight  curbw
eight \
0      1      3      88.6      168.8      64.1      48.8
2548
1      2      3      88.6      168.8      64.1      48.8
2548
2      3      1      94.5      171.2      65.5      52.4
2823
3      4      2      99.8      176.6      66.2      54.3
2337
4      5      2      99.4      176.6      66.4      54.3
2824
..      ...      ...      ...      ...      ...      ...
...
200    201      -1      109.1      188.8      68.9      55.5
2952
201    202      -1      109.1      188.8      68.8      55.5
3049
202    203      -1      109.1      188.8      68.9      55.5
3012
203    204      -1      109.1      188.8      68.9      55.5
3217
204    205      -1      109.1      188.8      68.9      55.5
3062

```

```

      enginesize  boreratio  stroke  ...  cylindernumber_three  \
0      130      3.47      2.68  ...      0
1      130      3.47      2.68  ...      0
2      152      2.68      3.47  ...      0
3      109      3.19      3.40  ...      0
4      136      3.19      3.40  ...      0
..      ...      ...      ...  ...      ...
200    141      3.78      3.15  ...      0
201    141      3.78      3.15  ...      0
202    173      3.58      2.87  ...      0
203    145      3.01      3.40  ...      0
204    141      3.78      3.15  ...      0

```

```

      cylindernumber_twelve  cylindernumber_two  fuelsystem_2bbl  \
0      0      0      0
1      0      0      0
2      0      0      0
3      0      0      0
4      0      0      0
..      ...      ...      ...
200    0      0      0
201    0      0      0
202    0      0      0
203    0      0      0
204    0      0      0

```

```

      fuelsystem_4bbl  fuelsystem_idi  fuelsystem_mfi  fuelsystem_mphi  \
0      0      0      0      1
1      0      0      0      1
2      0      0      0      1
3      0      0      0      1

```



4	0	0	0	1
..	...	...	...	...
200	0	0	0	1
201	0	0	0	1
202	0	0	0	1
203	0	1	0	0
204	0	0	0	1

	fuelsystem_spdi	fuelsystem_spfi
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
..	...	...
200	0	0
201	0	0
202	0	0
203	0	0
204	0	0

[205 rows x 190 columns]

0	13495.0
1	16500.0
2	16500.0
3	13950.0
4	17450.0

..	...
200	16845.0
201	19045.0
202	21485.0
203	22470.0
204	22625.0

Name: price, Length: 205, dtype: float64  
(205, 190)  
(205,)

```
In [11]: ## Training and Test Data ##
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=1/10,rand
```

```
In [12]: ## Linear Regression ##

from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train,y_train)
```

Out[12]: LinearRegression()

```
In [24]: from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(x_train, y_train)
```

Out[24]: LinearRegression()

```
In [25]: t_data_predic = regressor.predict(x_train)
```


```
In [26]: ## Error Calculation ##

from sklearn import metrics
error_score = metrics.r2_score(y_train, t_data_predic)
print("R squared Error : ", error_score)
```

R squared Error : 0.9982334777472928

```
In [28]: ## prediction on Test data ##
t_data_predic = regressor.predict(x_test)
```

```
In [30]: plt.scatter(y_test, t_data_predic)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")
plt.show()
```



```
In [31]: ▶ plt.scatter(y_test, t_data_predic)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")
plt.plot([min(y_train), max(y_train)], [min(y_train), max(y_train)], 'r--')
plt.show()
```



```
In [32]: b = regressor.coef_
print("Coefficient :",b)
```

```
Coefficient : [ 1.96509785e+02  6.44936604e+02  6.43751612e+03 -1.846823
73e+02
 9.62864785e+02  3.51905038e+02  1.47578359e+02 -2.76898277e+02
 1.30471318e+03 -4.90883212e+02  1.01640031e+01  2.07277966e+01
-2.61585478e+03 -1.46815247e+03 -9.51029627e+02 -2.58646302e+00
 2.42849538e+00  4.72045187e+02 -2.83591892e+02  2.23256072e+04
 1.60797850e+04  1.88882752e+04  2.16973985e+04  7.65695439e+03
 1.02018548e+04 -2.04272510e-09  5.52651730e+03 -3.35398909e-09
 2.37820393e+04  6.77573553e-10  3.50549631e+04  3.68539059e+04
 4.56932926e+04  3.43660050e+04  1.89276813e+03  2.29997800e+03
 7.59529123e+03  1.14404830e+02  2.14823528e+03  1.22054853e+04
 3.46955061e+03 -4.04518514e+03 -3.54134500e-11  9.33059554e+03
 9.01911563e+03  9.92581143e+03  7.63789167e+03  6.91740145e+03
 9.24904326e+03  4.63287726e+03  1.06968981e+04  1.42257065e+04
 7.39882614e+03  7.34533606e+03  2.02150903e+04  8.09450285e-10
 1.48353204e+04  1.24647316e+04  6.49764458e+03  1.56359652e+04
 1.57687221e+04  1.22894361e+04  1.61300267e+04  6.51504408e+03
 1.38243195e-10  1.06638524e+04  2.19892515e+04  1.88857612e+04
 2.13645926e+03  6.23869197e+03  4.12374372e+03  5.87180380e+03
 1.24491438e+04  6.07284177e+03  6.44050795e+03  4.32427242e+03
 7.00253541e+03  6.15724511e+03  1.40018140e+04 -3.27418093e-11
 5.56319813e+03 -2.10148582e+03 -3.23059905e+03  3.29237082e-10
-2.72483607e+03 -3.98786800e+03 -3.48772963e+03 -3.39093791e+03
-2.16371492e+03 -1.64521400e+03 -6.96119570e+02 -1.02350858e+03
-2.73538678e+02 -3.50329927e+03 -9.36155914e+02  1.87282882e+02
 3.21383449e+02 -1.99899510e+03 -9.72803154e+02 -8.19781102e+02
-7.49423634e-10  5.89554629e+02 -3.24370948e+03 -1.10929919e+03
-2.15366884e+02  9.72617050e+01 -1.24322448e+03 -9.65972082e+03
-1.08984822e+04 -1.05788393e+04 -1.18093123e+04 -9.08037847e+03
-4.41656453e+03  5.61981388e+03  7.80571692e+03  6.92330409e+03
-1.79391100e+03 -1.52837459e-09 -7.63639974e+03 -3.74613207e+03
-2.66239910e+03 -5.25487633e+03 -5.66331032e+03 -2.34515357e+03
-3.20881020e+03 -3.84748368e+03 -5.39207732e+03 -2.54882668e+03
-5.36876918e+03 -2.18861445e+03 -1.77123462e+04 -1.57492574e+04
-1.87165226e+04 -1.40786664e+04 -1.44430878e+04  4.69055864e-10
-1.42069626e+04 -1.85044657e+04 -1.46322479e+04 -1.08615203e+04
-1.51889079e+04 -1.56209964e+04 -1.23557763e+04 -1.42267862e+04
-1.23749457e+04 -8.20364221e-10 -1.36697054e-09 -1.80671048e+04
-1.74424691e+04 -1.91456875e+04 -1.72969657e+04 -1.92334812e+04
-2.83890813e+04 -1.65182732e+04 -1.90617839e+04 -1.86088455e+04
-1.69993612e+04 -1.79981202e+04 -1.51780806e+04 -1.50203418e+04
-1.19135692e+04 -1.37402866e+04 -1.35996658e+04 -1.97564832e+04
-1.95290640e+04 -1.73213477e+03 -1.20782794e+03 -2.20250857e+03
-1.47386125e+03  2.03488349e+04  0.00000000e+00 -5.71433833e+03
-2.69459650e+03 -1.02142105e+04  2.12168539e+03 -2.84346726e+03
 2.77294682e+03 -6.79229506e+03 -1.31774704e+04 -1.70530257e-13
 2.13645926e+03 -2.84346726e+03  2.12654196e+03 -5.77284403e+03
 4.63287726e+03  6.43751612e+03 -8.91630903e+03  1.49167956e+03
 1.27849341e+03]
```

```
In [33]: ▶ a = regressor.intercept_  
print("Intercept :",a)
```

Intercept : -45068.74680415792