

```
In [1]: import pandas as pd
import numpy as np
```


```
In [2]: df = pd.read_csv('CarPrice.csv')
```

```
In [3]: df
```

```
Out[3]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engir
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	
...
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rwd	
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rwd	
202	203	-1	volvo 244dl	gas	std	four	sedan	rwd	
203	204	-1	volvo 246	diesel	turbo	four	sedan	rwd	
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rwd	

205 rows × 26 columns




```
In [4]: ## DATA PREPROCESSING ##
```

```
In [4]: df.head()
```

Out[4]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	

5 rows × 26 columns




In [5]: `df.tail()`

Out[5]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	engine
200	201	-1	volvo 145e (sw)	gas	std	four	sedan	rwd	
201	202	-1	volvo 144ea	gas	turbo	four	sedan	rwd	
202	203	-1	volvo 244dl	gas	std	four	sedan	rwd	
203	204	-1	volvo 246	diesel	turbo	four	sedan	rwd	
204	205	-1	volvo 264gl	gas	turbo	four	sedan	rwd	

5 rows × 26 columns



In [6]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   car_ID                205 non-null   int64
 1   symboling              205 non-null   int64
 2   CarName                205 non-null   object
 3   fueltype              205 non-null   object
 4   aspiration             205 non-null   object
 5   doornumber            205 non-null   object
 6   carbody               205 non-null   object
 7   drivewheel            205 non-null   object
 8   enginelocation        205 non-null   object
 9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm               205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg            205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB

```

```
In [7]: df.shape
```

```
Out[7]: (205, 26)
```

```
In [8]: df.columns
```

```
Out[8]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
             'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
             'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
             'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
             'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
             'price'],
            dtype='object')
```

```
In [9]: df.duplicated().sum()
```

```
Out[9]: 0
```

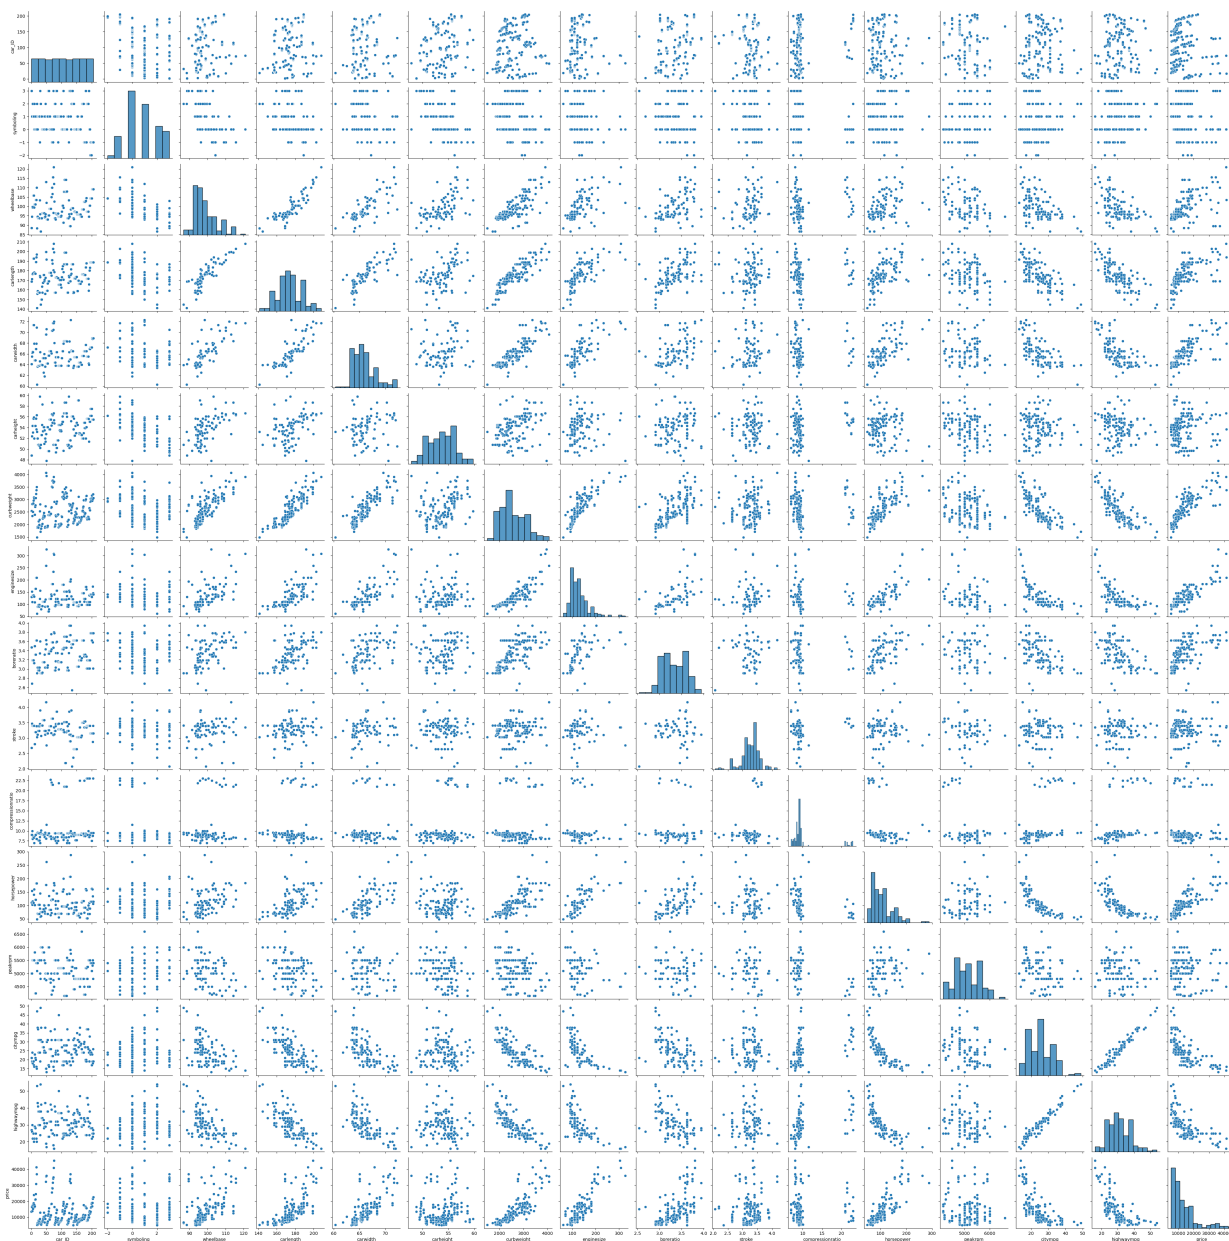
```
In [10]: df.isnull().sum()
```

```
Out[10]: car_ID          0
          symboling     0
          CarName        0
          fueltype       0
          aspiration     0
          doornumber     0
          carbody        0
          drivewheel     0
          enginelocation 0
          wheelbase      0
          carlength      0
          carwidth       0
          carheight      0
          curbweight     0
          enginetype     0
          cylindernumber 0
          enginesize     0
          fuelsystem     0
          boreratio      0
          stroke         0
          compressionratio 0
          horsepower     0
          peakrpm        0
          citympg        0
          highwaympg     0
          price          0
          dtype: int64
```

```
In [11]: import matplotlib.pyplot as plt
          import seaborn as sns

          sns.pairplot(df)
```

```
Out[11]: <seaborn.axisgrid.PairGrid at 0x1f5eb6eef70>
```



In [12]: **## CLASSIFICATION OF CATEGORICAL DATA ##**

```
print(df.fueltype.value_counts())
print(df.aspiration.value_counts())
print(df.doornumber.value_counts())
print(df.carbody.value_counts())
print(df.drivewheel.value_counts())
print(df.fuelsystem.value_counts())
print(df.cylindernumber.value_counts())
```

```
gas      185
diesel   20
Name: fueltype, dtype: int64
std      168
turbo    37
Name: aspiration, dtype: int64
four     115
two       90
Name: doornumber, dtype: int64
sedan     96
hatchback 70
wagon     25
hardtop    8
convertible 6
Name: carbody, dtype: int64
fwd      120
rwd       76
4wd        9
Name: drivewheel, dtype: int64
mpfi      94
2bbl      66
idi       20
1bbl      11
spdi       9
4bbl       3
mfi        1
spfi       1
Name: fuelsystem, dtype: int64
four      159
six        24
five       11
eight       5
two         4
three       1
twelve      1
Name: cylindernumber, dtype: int64
```

```
In [13]: df['horsepower'].value_counts()
```

```
Out[13]: 68      19
          70      11
          69      10
          116      9
          110      8
          95      7
          114      6
          160      6
          101      6
          62      6
          88      6
          145      5
          76      5
          97      5
          84      5
          90      5
          82      5
          102      5
          92      4
          111      4
          123      4
          86      4
          207      3
          73      3
          182      3
          121      3
          85      3
          152      3
          176      2
          94      2
          56      2
          112      2
          161      2
          184      2
          155      2
          156      2
          52      2
          100      2
          162      2
          140      1
          115      1
          134      1
          78      1
          142      1
          288      1
          143      1
          48      1
          200      1
          58      1
          55      1
          60      1
          175      1
          154      1
          72      1
          120      1
          64      1
          135      1
          262      1
          106      1
```

Name: horsepower, dtype: int64

```
In [14]: df['stroke'].value_counts()
```

```
Out[14]: 3.400    20
          3.230    14
          3.150    14
          3.030    14
          3.390    13
          2.640    11
          3.290     9
          3.350     9
          3.460     8
          3.110     6
          3.270     6
          3.410     6
          3.070     6
          3.580     6
          3.190     6
          3.500     6
          3.640     5
          3.520     5
          3.860     4
          3.540     4
          3.470     4
          3.255     4
          3.900     3
          2.900     3
          3.100     2
          4.170     2
          2.800     2
          2.190     2
          3.080     2
          2.680     2
          2.360     1
          3.160     1
          2.070     1
          3.210     1
          3.120     1
          2.760     1
          2.870     1
          Name: stroke, dtype: int64
```

```
In [15]: df['compressionratio'].value_counts()
```



```
Out[15]:
```

9.00	46
9.40	26
8.50	14
9.50	13
9.30	11
8.70	9
8.00	8
9.20	8
7.00	7
8.60	5
21.00	5
8.40	5
7.50	5
23.00	5
9.60	5
21.50	4
7.60	4
10.00	3
22.50	3
8.30	3
8.80	3
7.70	2
8.10	2
9.10	1
9.31	1
7.80	1
9.41	1
21.90	1
22.00	1
22.70	1
10.10	1
11.50	1

Name: compressionratio, dtype: int64

```
In [16]: df['citympg'].value_counts()
```

```
Out[16]: 31    28
          19    27
          24    22
          27    14
          17    13
          26    12
          23    12
          21     8
          25     8
          30     8
          38     7
          28     7
          16     6
          37     6
          22     4
          29     3
          15     3
          20     3
          18     3
          14     2
          34     1
          35     1
          32     1
          36     1
          45     1
          13     1
          49     1
          47     1
          33     1
          Name: citympg, dtype: int64
```

```
In [17]: df['highwaympg'].value_counts()
```

```
Out[17]: 25    19
          38    17
          24    17
          30    16
          32    16
          34    14
          37    13
          28    13
          29    10
          33     9
          22     8
          31     8
          23     7
          27     5
          43     4
          42     3
          26     3
          41     3
          19     2
          39     2
          18     2
          16     2
          20     2
          36     2
          47     2
          46     2
          54     1
          17     1
          53     1
          50     1
Name: highwaympg, dtype: int64
```

```
In [18]: ## CHANGING THE CATEGORICAL ATTRIBUTES INTO NUMERIC DATA FOR BETTER ANALYSIS ##

df.replace({'fueltype':{'gas':0,'diesel':1}},inplace=True)
df.replace({'aspiration':{'std':0,'turbo':1}},inplace=True)
df.replace({'doornumber':{'two':0,'four':1}},inplace=True)
df.replace({'carbody':{'convertible':0,'hatchback':1,'sedan':2, 'wagon':3}},inplace=True)
df.replace({'drivewheel':{'rwd':0,'fwd':1,'4wd':2}},inplace=True)
df.replace({'fuelsystem':{'mpfi':0,'2bbl':1,'1bbl':2,'mfi':3, 'spf1':4, 'idi':5}},inplace=True)
```

```
In [19]: ## hot encoding ##
df = pd.get_dummies(df,drop_first=True)
df
```

Out[19]:

	car_ID	symboling	fueltype	aspiration	doornumber	drivewheel	wheelbase	carlength	carwidth
0	1	3	0	0	0	0	88.6	168.8	64
1	2	3	0	0	0	0	88.6	168.8	64
2	3	1	0	0	0	0	94.5	171.2	65
3	4	2	0	0	1	1	99.8	176.6	66
4	5	2	0	0	1	2	99.4	176.6	66
...
200	201	-1	0	0	1	0	109.1	188.8	68
201	202	-1	0	1	1	0	109.1	188.8	68
202	203	-1	0	0	1	0	109.1	188.8	68
203	204	-1	1	1	1	0	109.1	188.8	68
204	205	-1	0	1	1	0	109.1	188.8	68

205 rows × 190 columns

In [20]: *## Splitting the data ##*

```

x = df.drop(['price'], axis=1)
y = df['price']
print(len(x), len(y))
print(x)
print(y)
print(x.shape)
print(y.shape)

```

205 205

	car_ID	symboling	fueltype	aspiration	doornumber	drivewheel	\
0	1	3	0	0	0	0	
1	2	3	0	0	0	0	
2	3	1	0	0	0	0	
3	4	2	0	0	1	1	
4	5	2	0	0	1	2	
..	
200	201	-1	0	0	1	0	
201	202	-1	0	1	1	0	
202	203	-1	0	0	1	0	
203	204	-1	1	1	1	0	
204	205	-1	0	1	1	0	

	wheelbase	carlength	carwidth	carheight	...	cylindernumber_three	\
0	88.6	168.8	64.1	48.8	...	0	
1	88.6	168.8	64.1	48.8	...	0	
2	94.5	171.2	65.5	52.4	...	0	
3	99.8	176.6	66.2	54.3	...	0	
4	99.4	176.6	66.4	54.3	...	0	
..	
200	109.1	188.8	68.9	55.5	...	0	
201	109.1	188.8	68.8	55.5	...	0	
202	109.1	188.8	68.9	55.5	...	0	
203	109.1	188.8	68.9	55.5	...	0	
204	109.1	188.8	68.9	55.5	...	0	

	cylindernumber_twelve	cylindernumber_two	fuelsystem_1	fuelsystem_2	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
..	
200	0	0	0	0	
201	0	0	0	0	
202	0	0	0	0	
203	0	0	0	0	
204	0	0	0	0	

	fuelsystem_3	fuelsystem_5	fuelsystem_4bbl	fuelsystem_spdi	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
..	
200	0	0	0	0	
201	0	0	0	0	
202	0	0	0	0	
203	0	1	0	0	
204	0	0	0	0	

	fuelsystem_spfi
0	0
1	0
2	0
3	0
4	0
..	...

```

200      0
201      0
202      0
203      0
204      0

```

```
[205 rows x 189 columns]
```

```

0      13495.0
1      16500.0
2      16500.0
3      13950.0
4      17450.0

```

```
...
```

```

200      16845.0
201      19045.0
202      21485.0
203      22470.0
204      22625.0

```

```

Name: price, Length: 205, dtype: float64
(205, 189)
(205,)

```

```

In [21]: ## Training and Test Data ##
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=1/10,random_state=0)

```

```

In [22]: ## Linear Regression ##

from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train,y_train)
model.score(x_test,y_test)

```

```
Out[22]: -2.140565648522962
```

```

In [23]: from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(x_train, y_train)

```

```
Out[23]: LinearRegression()
```

```
In [24]: t_data_predic = regressor.predict(x_train)
```

```

In [25]: ## Error Calculation ##

from sklearn import metrics
error_score = metrics.r2_score(y_train, t_data_predic)
print("R squared Error : ", error_score)

```


```
R squared Error : 0.9982334777472928
```

```

In [26]: ## Plotting THE data ##

plt.scatter(y_train, t_data_predic)
plt.xlabel("Actual Price of Car")
plt.ylabel("Predicted Price of Car")
plt.title(" Actual Prices vs Predicted Prices")
plt.show()

```



```
In [31]: ## prediction on Training data ##  
t_data_predic = regressor.predict(x_test)
```

```
In [32]: # R squared Error ##  
error_score = metrics.r2_score(y_test, t_data_predic)  
print("R squared Error : ", error_score)
```

R squared Error : -2.140565648522962

```
In [33]: plt.scatter(y_test, t_data_predic)  
plt.xlabel("Actual Price")  
plt.ylabel("Predicted Price")  
plt.title(" Actual Prices vs Predicted Prices")  
plt.show()
```



In []:

In []: